



Operating Systems

System Calls

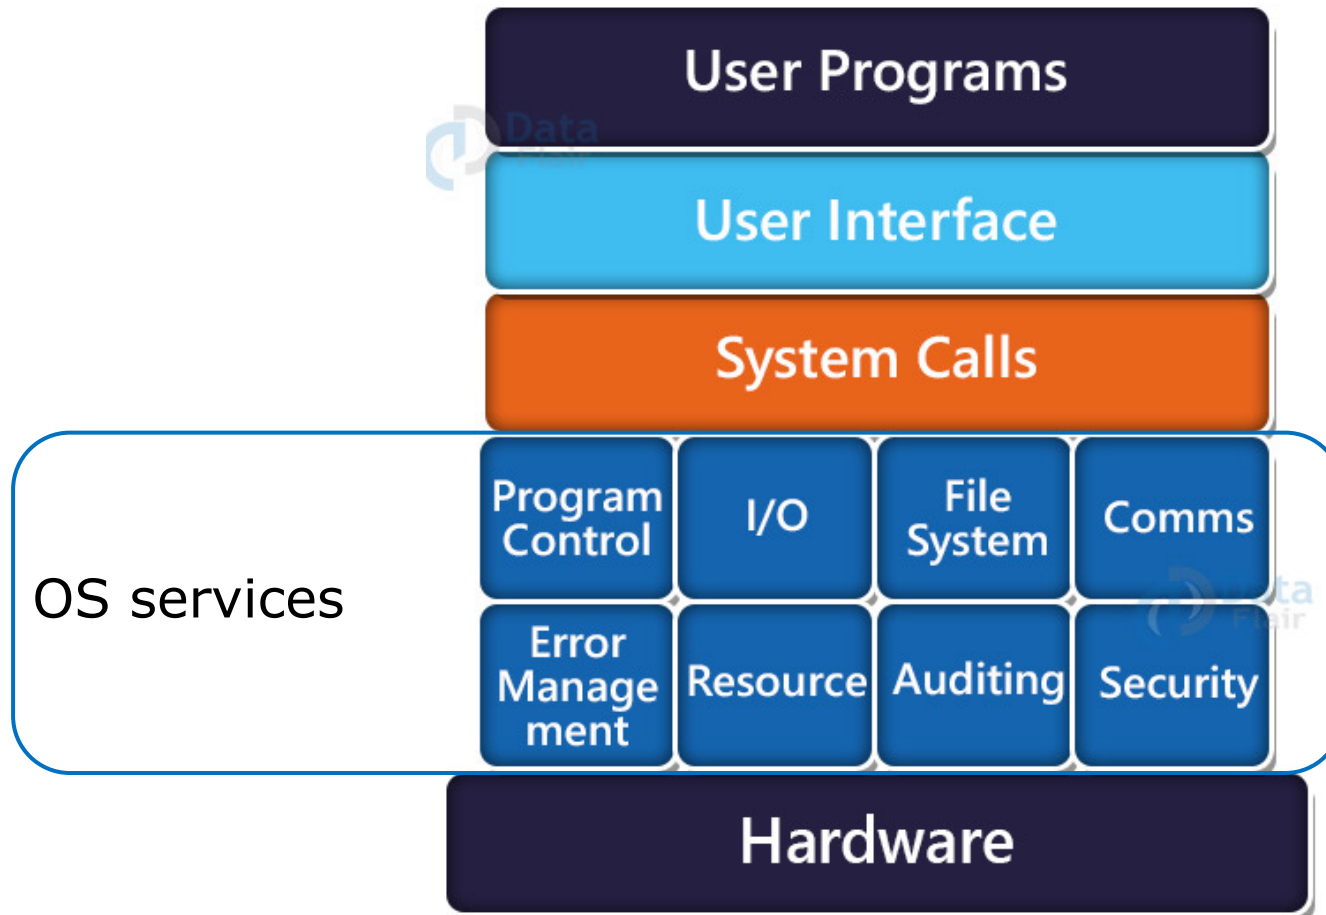
Seyyed Ahmad Javadi

sajavadi@aut.ac.ir

Spring 2023

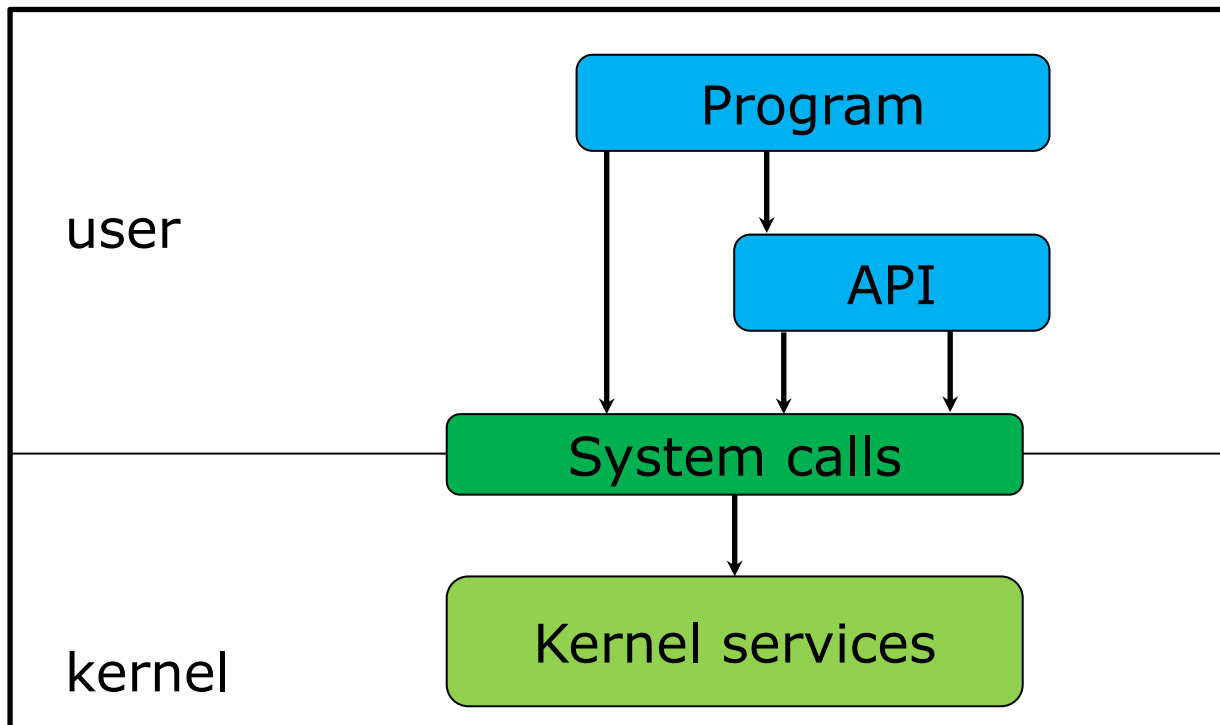
System Calls

- **Programming interface** to the **services** provided by the OS.



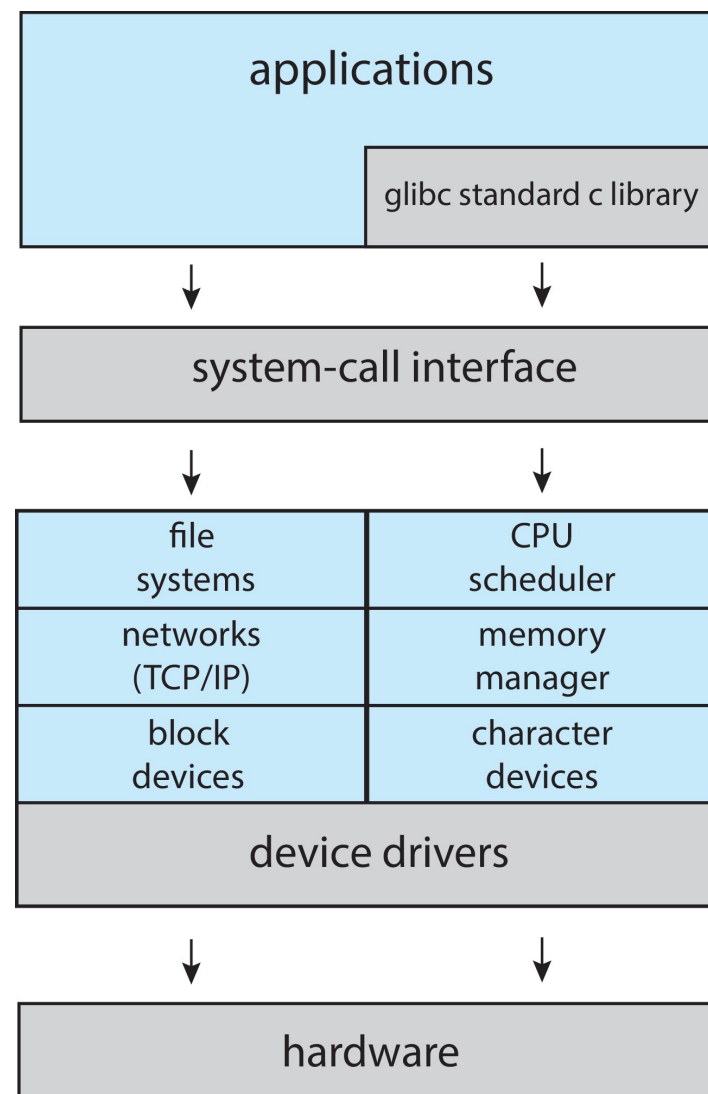
System Calls (cont.)

- Typically written in a high-level language (C or C++ or Assembly).
- Mostly accessed by programs via a high-level **Application Programming Interface (API)** rather than direct system call use.



Three most common APIs

- Win32 API for Windows (Win API)
- POSIX API for POSIX-based systems
 - Including virtually all versions of UNIX, Linux (*unistd.h*), and Mac OS X
- Java API for the Java virtual machine.



Example of Standard API

EXAMPLE OF STANDARD API

As an example of a standard API, consider the `read()` function that is available in UNIX and Linux systems. The API for this function is obtained from the man page by invoking the command

```
man read
```

on the command line. A description of this API appears below:

```
#include <unistd.h>
```

```
ssize_t
```

```
read(int fd, void *buf, size_t count)
```

return
value

function
name

parameters

Example of Standard API (cont.)

A program that uses the `read()` function must include the `unistd.h` header file, as this file defines the `ssize_t` and `size_t` data types (among other things). The parameters passed to `read()` are as follows:

- `int fd`—the file descriptor to be read
- `void *buf`—a buffer into which the data will be read
- `size_t count`—the maximum number of bytes to be read into the buffer

On a successful read, the number of bytes read is returned. A return value of 0 indicates end of file. If an error occurs, `read()` returns `-1`.



System Call Implementation

- Typically, a number is associated with each system call
 - **System-call interface** maintains a table indexed according to these numbers.
- The system call interface invokes the intended system call in OS kernel and **returns status of the system call** and any **return values**

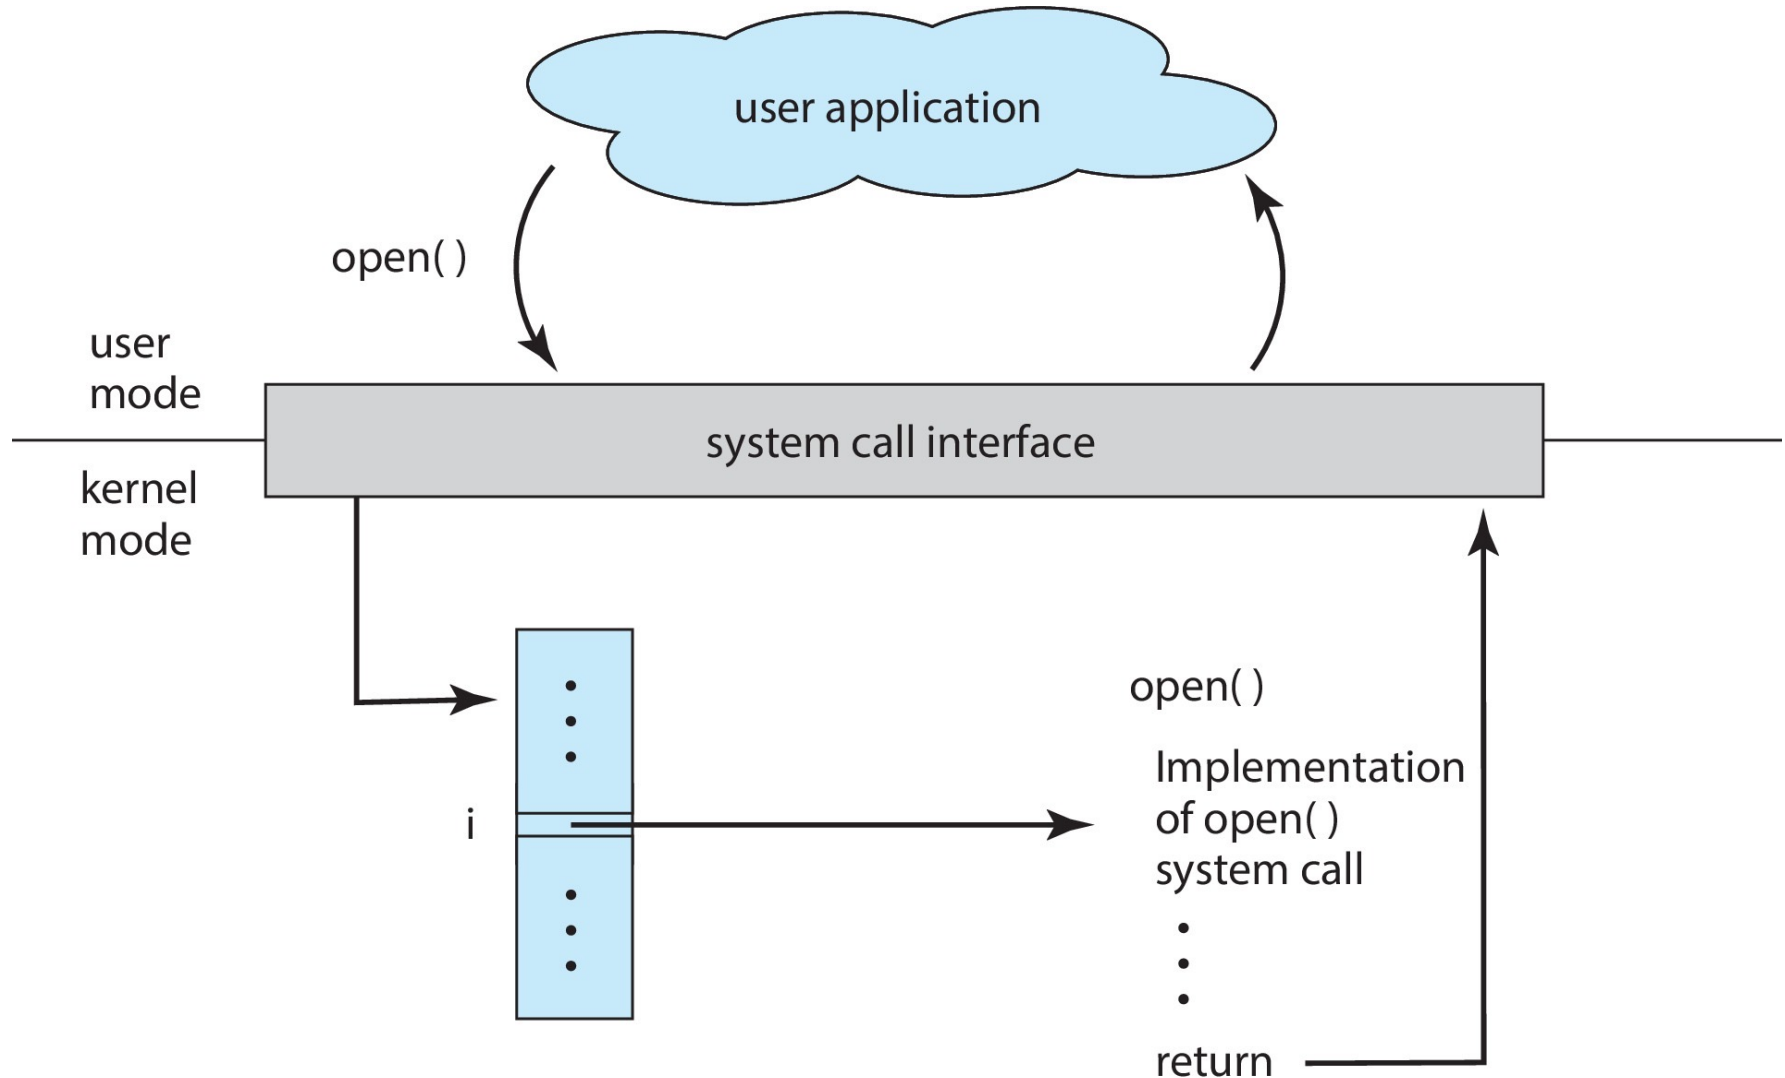


System Call Implementation (cont.)

- The caller need know nothing about how the system call is implemented
 - Just needs to obey API and understand what OS will do as a result call.
 - Most details of OS interface hidden from programmer by API
 - ▶ Managed by run-time support library



API – System Call – OS Relationship



System calls in assembly programs (demo)

- Put the system call number in the **EAX register**.
- Store the arguments to the system call in the registers EBX, ECX,...
- Call the relevant interrupt (**80h**).
- The result is **usually** returned in the **EAX** register.

https://www.tutorialspoint.com/assembly_programming/assembly_system_calls.htm

http://faculty.nps.edu/cseagle/assembly/sys_call.html



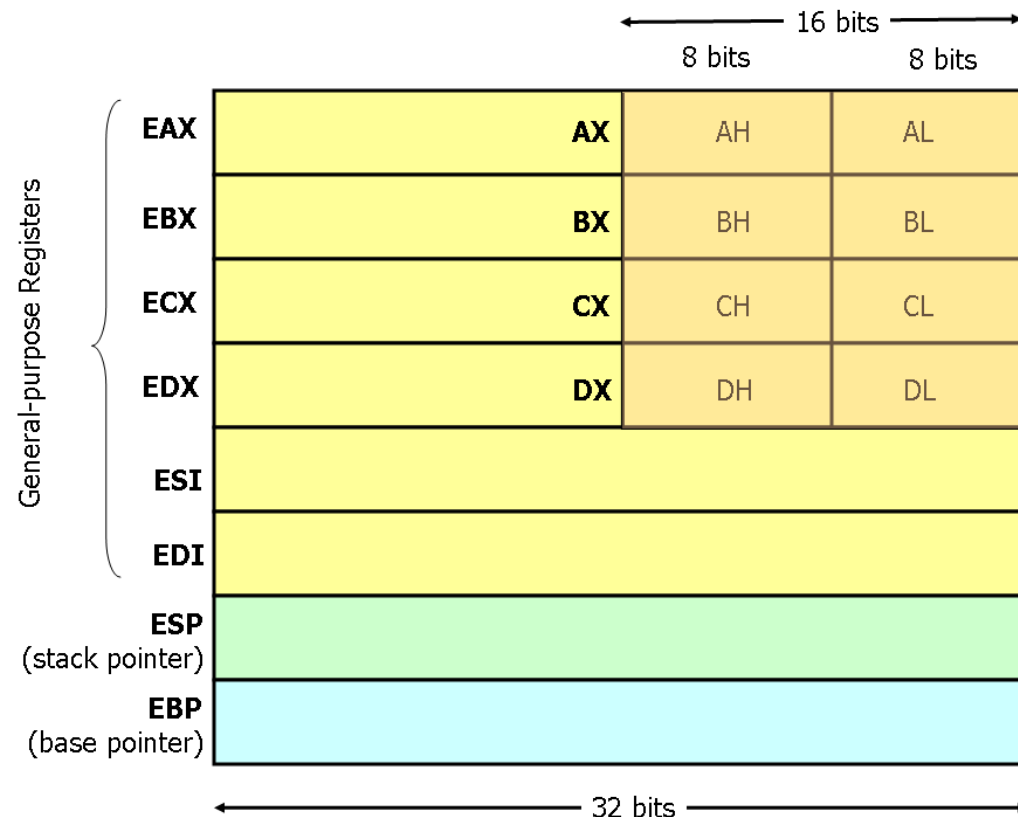
System Call Parameter Passing

- **Parameter Passing**
 - Register
 - Register pointer to mem. block
 - Stack (Push, Pop)
- Often, more information is required than simply identity of desired system call.
- Exact type and amount of information vary according to OS and call.



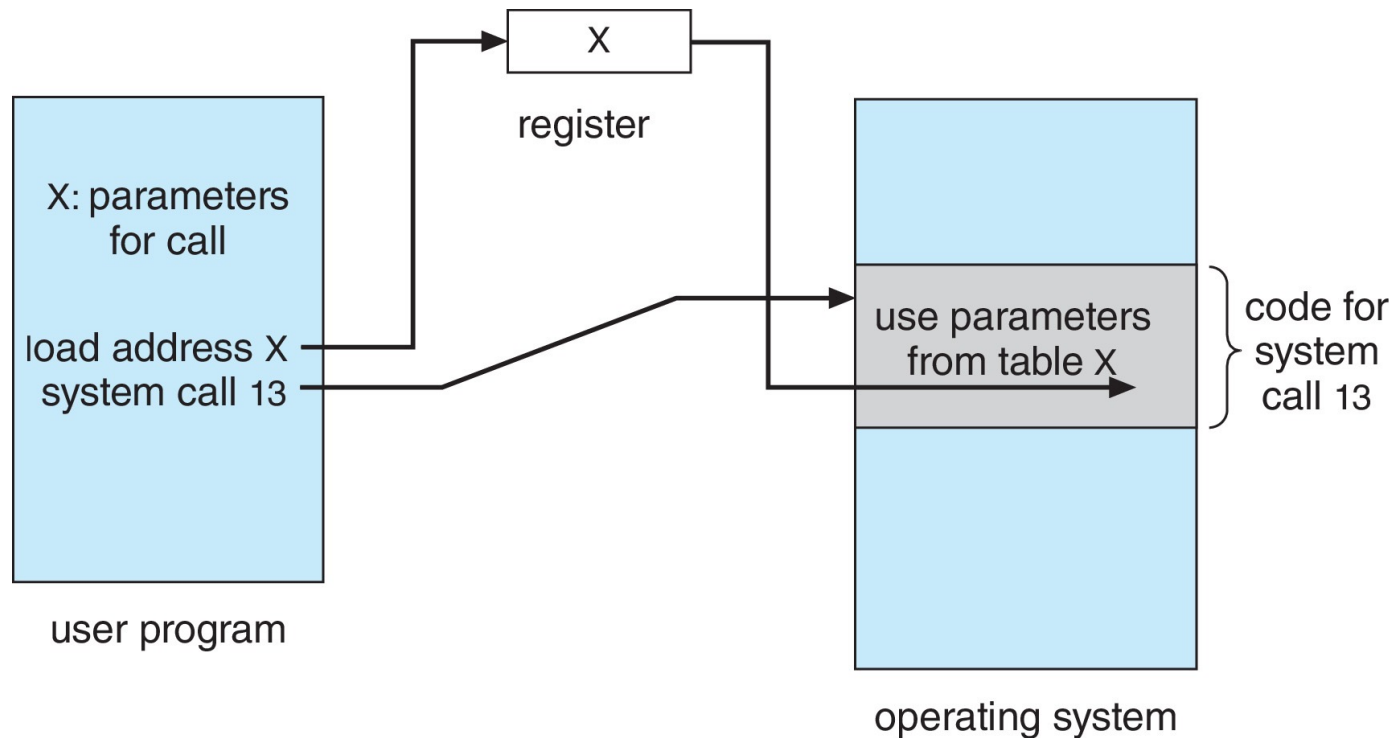
System Call Parameter Passing--Methods

- **Simplest:** pass the parameters in registers.
 - In some cases, may be more parameters than registers.




System Call Parameter Passing—Methods (cont.)

- Parameters stored in a **block, or table, in memory**, and address of block passed as a parameter in a register.
- This approach taken by Linux and Solaris.



System Call Parameter Passing—Methods (cont.)

- Parameters placed, or **pushed**, onto the **stack** by the program and **popped** off the stack by the operating system.



System Call Parameter Passing—Methods (cont.)

| Methods\features | Is there a limitation on the number of parameters? | Is there a limitation on the length of the parameters? |
|---------------------------------------|--|--|
| Register | | |
| Register pointer to mem. block | | |
| Stack (Push, Pop) | | |

System Call Parameter Passing—Methods (cont.)

| Methods\features | Is there a limitation on the number of parameters? | Is there a limitation on the length of the parameters? |
|---------------------------------------|--|--|
| Register | YES | YES |
| Register pointer to mem. block | NO | NO |
| Stack (Push, Pop) | NO | NO |



Types of System Calls

- **Process control**
 - Create process, terminate process
 - ...
- **File management**
 - create file, delete file
 - ...
- **Device management**
 - request device, release device
 - ...
- **Please study the reference book for more details**



Types of System Calls (Cont.)

| | Windows | Unix |
|--------------------------------|---|--|
| Process Control | CreateProcess() ExitProcess() WaitForSingleObject() | fork() exit() wait() |
| File Manipulation | CreateFile() ReadFile() WriteFile() CloseHandle() | open() read() write() close() |
| Device Manipulation | SetConsoleMode() ReadConsole() WriteConsole() | ioctl() read() write() |
| Information Maintenance | GetCurrentProcessID() SetTimer() Sleep() | getpid() alarm() sleep() |
| Communication | CreatePipe() CreateFileMapping() MapViewOfFile() | pipe() shm_open() mmap() |
| Protection | SetFileSecurity() InitializeSecurityDescriptor() SetSecurityDescriptorGroup() | chmod() umask() chown() |



Why Applications are Operating System Specific

- Apps compiled on one system usually not executable on other OSs.
- Each OS provides its own unique system calls
 - Own file formats, etc.
- Apps can be multi-operating system
 - Written in interpreted language like Python, Ruby, and interpreter available on multiple OSs.
 - App written in language that includes a VM containing the running app (like Java).
 - Use standard language (like C), compile separately on each operating system to run on each.

